

CMPT 412 Slides for Mean Shift


Edited down from longer Powerpoint by Yaron
Ukrainitz & Bernard Sarel that I retrieved
from the Internet

Mean Shift

Theory and Applications

Yaron Ukrainitz & Bernard Sarel

Brian's Note: You have to run this Powerpoint file in "Slide Show" mode. If you don't then the slide transitions don't work and lots of things get covered up.



The full original set of slides is available from
http://www.wisdom.weizmann.ac.il/~vision/courses/2004_2/files/mean_shift/mean_shift.ppt

Agenda

- **Mean Shift Theory**


- What is Mean Shift ?
- Density Estimation Methods
- Deriving the Mean Shift
- Mean shift properties

- **Applications**

- Clustering
- Discontinuity Preserving Smoothing
- Object Contour Detection
- Segmentation
- Object Tracking


Mean Shift Theory

Intuitive Description



Objective : Find the densest region
Distribution of identical billiard balls

Intuitive Description



Objective : Find the densest region
Distribution of identical billiard balls


Intuitive Description

Region of
interest


Center of
mass

Mean Shift
vector

Objective : Find the densest region
Distribution of identical billiard balls




Intuitive Description




Objective : Find the densest region
Distribution of identical billiard balls

Intuitive Description




Objective : Find the densest region
Distribution of identical billiard balls

Intuitive Description



Objective : Find the densest region
Distribution of identical billiard balls

Intuitive Description



Objective : Find the densest region
Distribution of identical billiard balls


What is Mean Shift ?

A tool for:

Finding modes in a set of data samples, manifesting an underlying probability density function (PDF) in \mathbb{R}^N


PDF in feature space

- Color space
- Scale space
- Actually any feature space you can conceive
- ...




Data

Non-parametric
Density **GRADIENT** Estimation
(Mean Shift)




PDF Representation



PDF Analysis

Non-Parametric Density Estimation


Assumption : The data points are sampled from an underlying PDF




Assumed Underlying PDF

Real Data Samples

Non-Parametric Density Estimation




Assumed Underlying PDF




Real Data Samples

Non-Parametric Density Estimation



Assumed Underlying PDF



Real Data Samples

Computing The Mean Shift

$$\nabla P(\mathbf{x}) = \frac{c}{n} \sum_{i=1}^n \nabla k_i = \frac{c}{n} \left[\sum_{i=1}^n g_i \right] g \left[\frac{\sum_{i=1}^n \mathbf{x}_i g_i}{\sum_{i=1}^n g_i} - \mathbf{x} \right]$$


Yet another Kernel density estimation !

Simple Mean Shift procedure:

- Compute mean shift vector


$$\mathbf{m}(\mathbf{x}) = \left[\frac{\sum_{i=1}^n \mathbf{x}_i g \left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h} \right)}{\sum_{i=1}^n g \left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h} \right)} - \mathbf{x} \right]$$

- Translate the Kernel window by $\mathbf{m}(\mathbf{x})$



$$g(\mathbf{x}) = -k'(\mathbf{x})$$


Mean Shift Mode Detection



Updated Mean Shift Procedure:


- Find all modes using the Simple Mean Shift Procedure
- Prune modes by perturbing them (find saddle points and plateaus)
- Prune nearby – take highest mode in the window

Mean Shift Properties



- Automatic convergence speed – the mean shift vector size depends on the gradient itself.
 - Near maxima, the steps are small and refined
 - Convergence is guaranteed for infinitesimal steps only → infinitely convergent, (therefore set a lower bound)
 - For Uniform Kernel (🌈), convergence is achieved in a finite number of steps
 - Normal Kernel (📐) exhibits a smooth trajectory, but is slower than Uniform Kernel (🌈).
- Adaptive**
Gradient
Ascent


Real Modality Analysis



Tessellate the space
with windows

Run the procedure in parallel


Real Modality Analysis



The blue data points were traversed by the windows towards the mode


Real Modality Analysis

An example



Window tracks signify the steepest ascent directions

Mean Shift Strengths & Weaknesses



Strengths :

- Application independent tool
- Suitable for real data analysis
- Does not assume any prior shape (e.g. elliptical) on data clusters
- Can handle arbitrary feature spaces
- Only ONE parameter to choose
- h (window size) has a physical meaning, unlike K-Means

Weaknesses :


- The window size (bandwidth selection) is not trivial
- Inappropriate window size can cause modes to be merged, or generate additional “shallow” modes → Use adaptive window size

Mean Shift Applications

Clustering



Cluster : All data points in the **attraction basin** of a mode

Attraction basin : the region for which all trajectories lead to the same mode





Clustering

Synthetic Examples



Simple Modal Structures




Complex Modal Structures

Clustering


Real Example

Feature space:
 L^*u^*v representation


Initial window
centers




(a)




(b)



Modes found (c)




Modes after
pruning (d)



Final clusters (e)

Clustering

Real Example




L*u*v space representation


Clustering

Real Example

2D (L^*u)
space
representation




(a)



(b)


Final clusters

Not all trajectories
in the attraction basin
reach the same mode



(c)

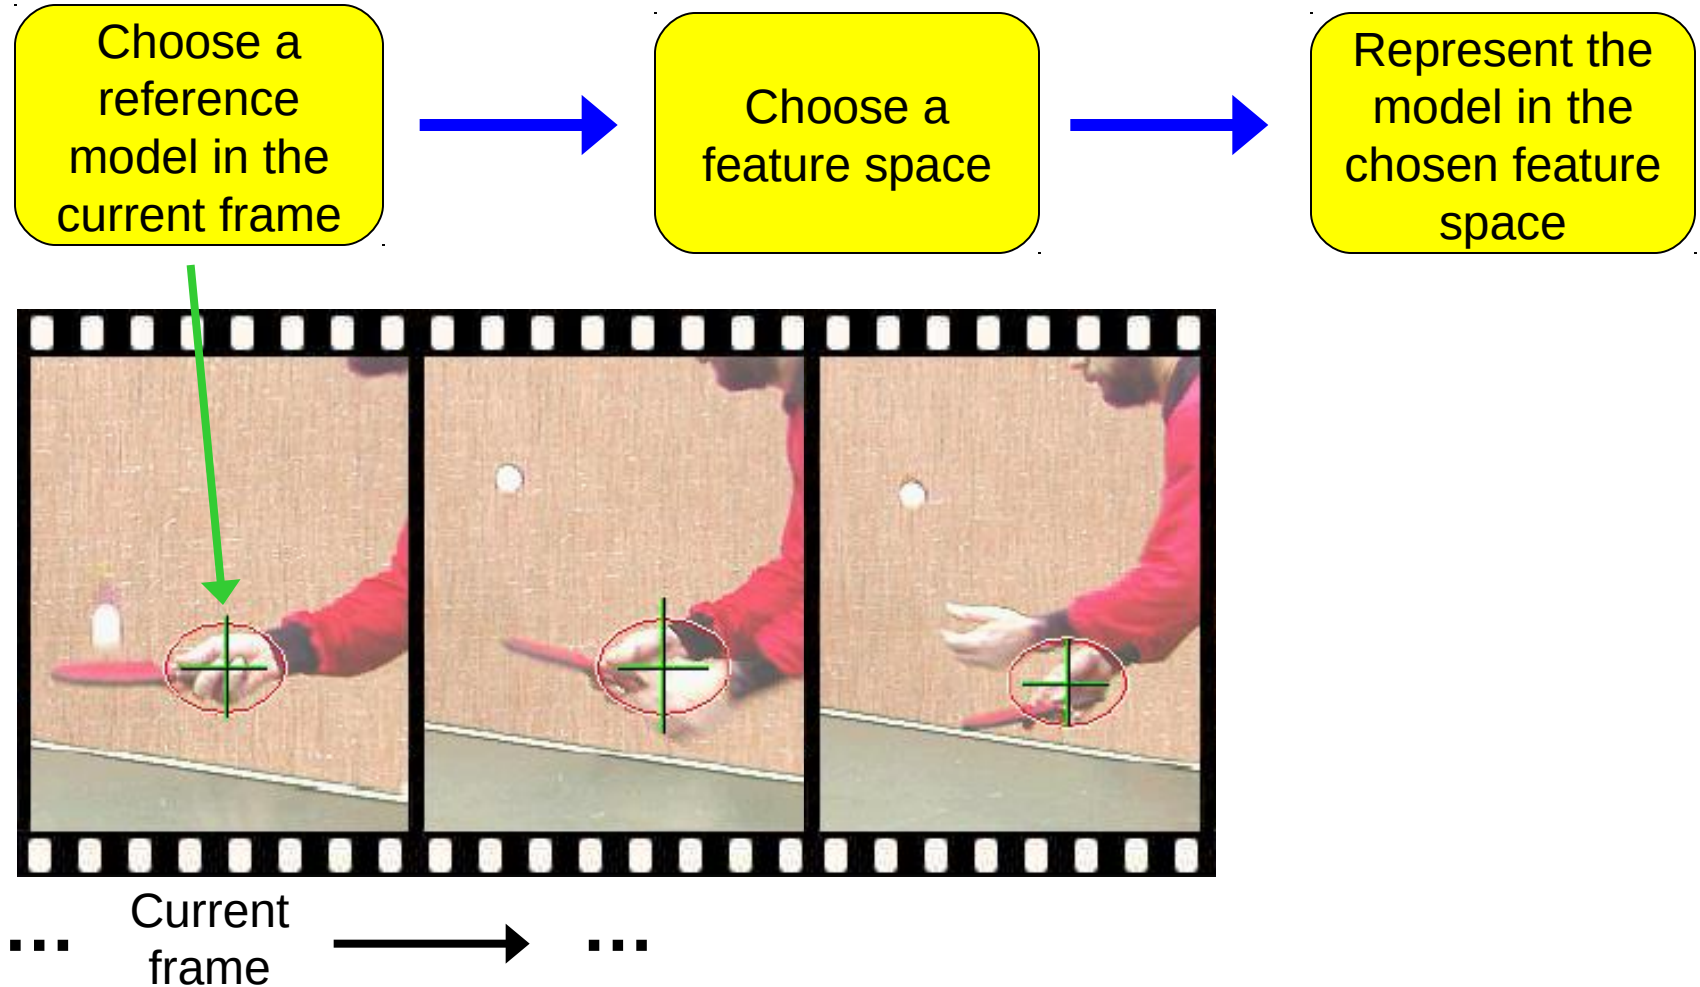
Non-Rigid Object Tracking



... → ...

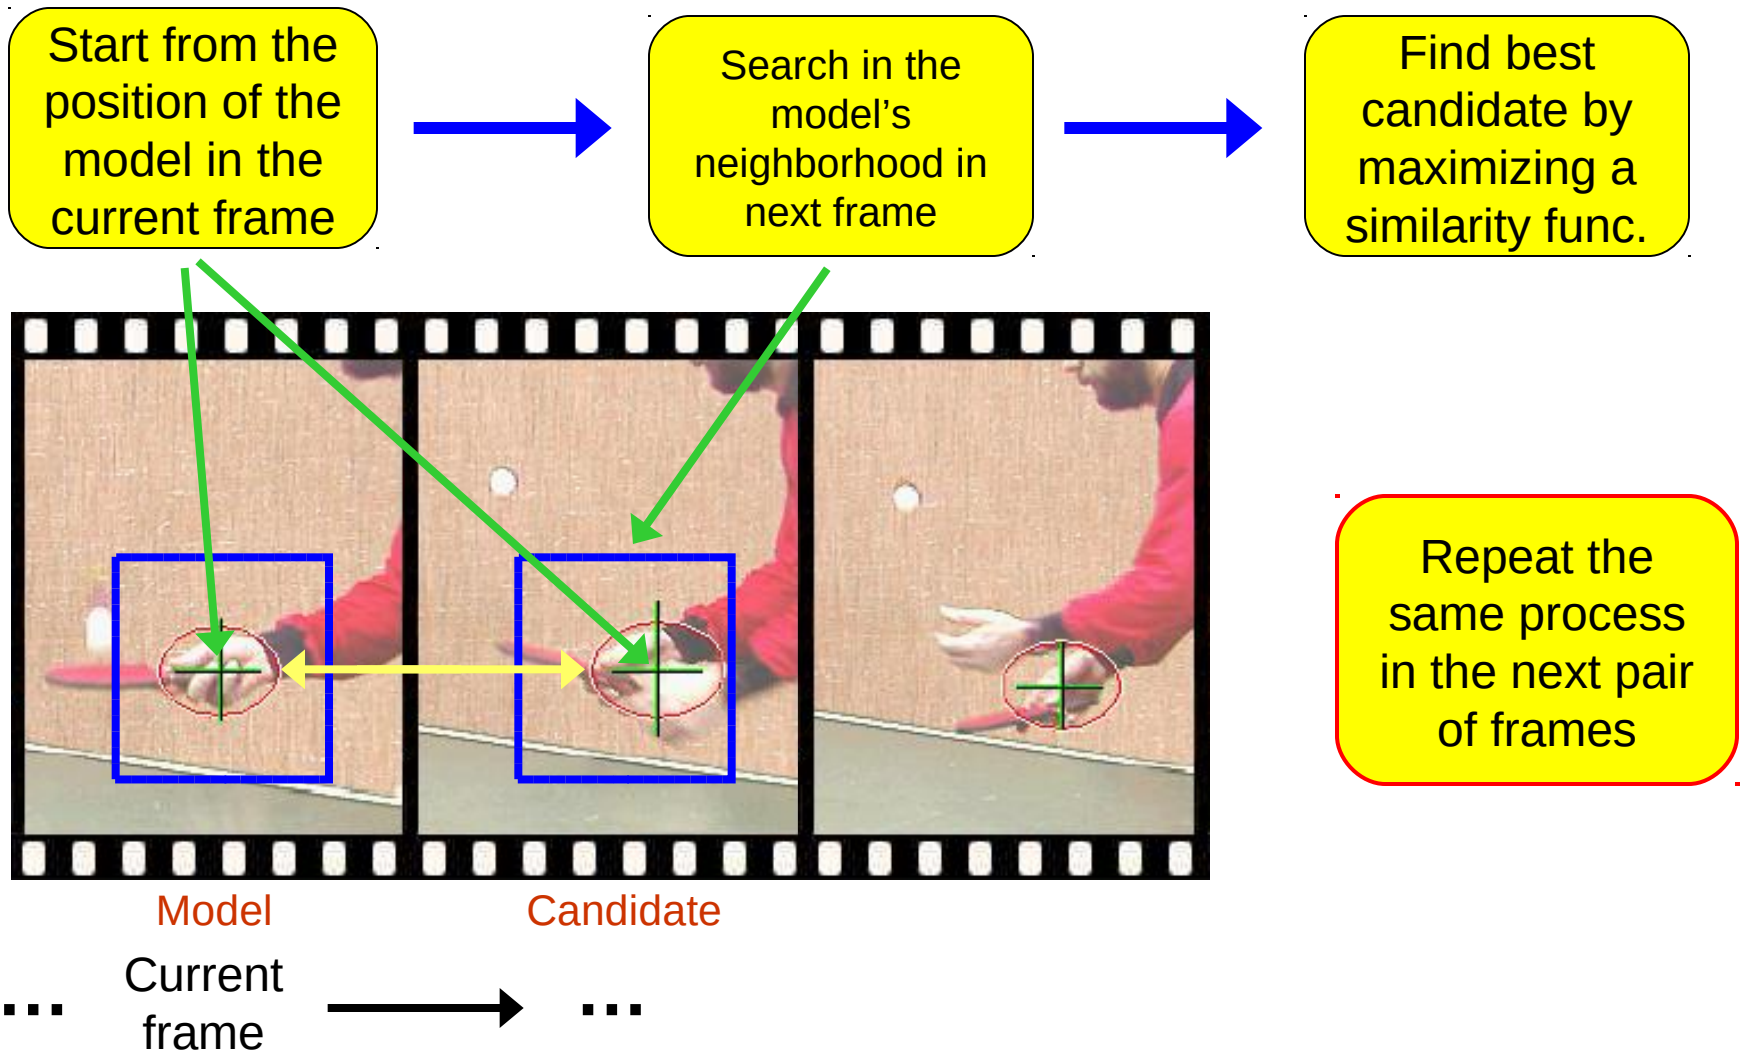
Mean-Shift Object Tracking

General Framework: Target Representation



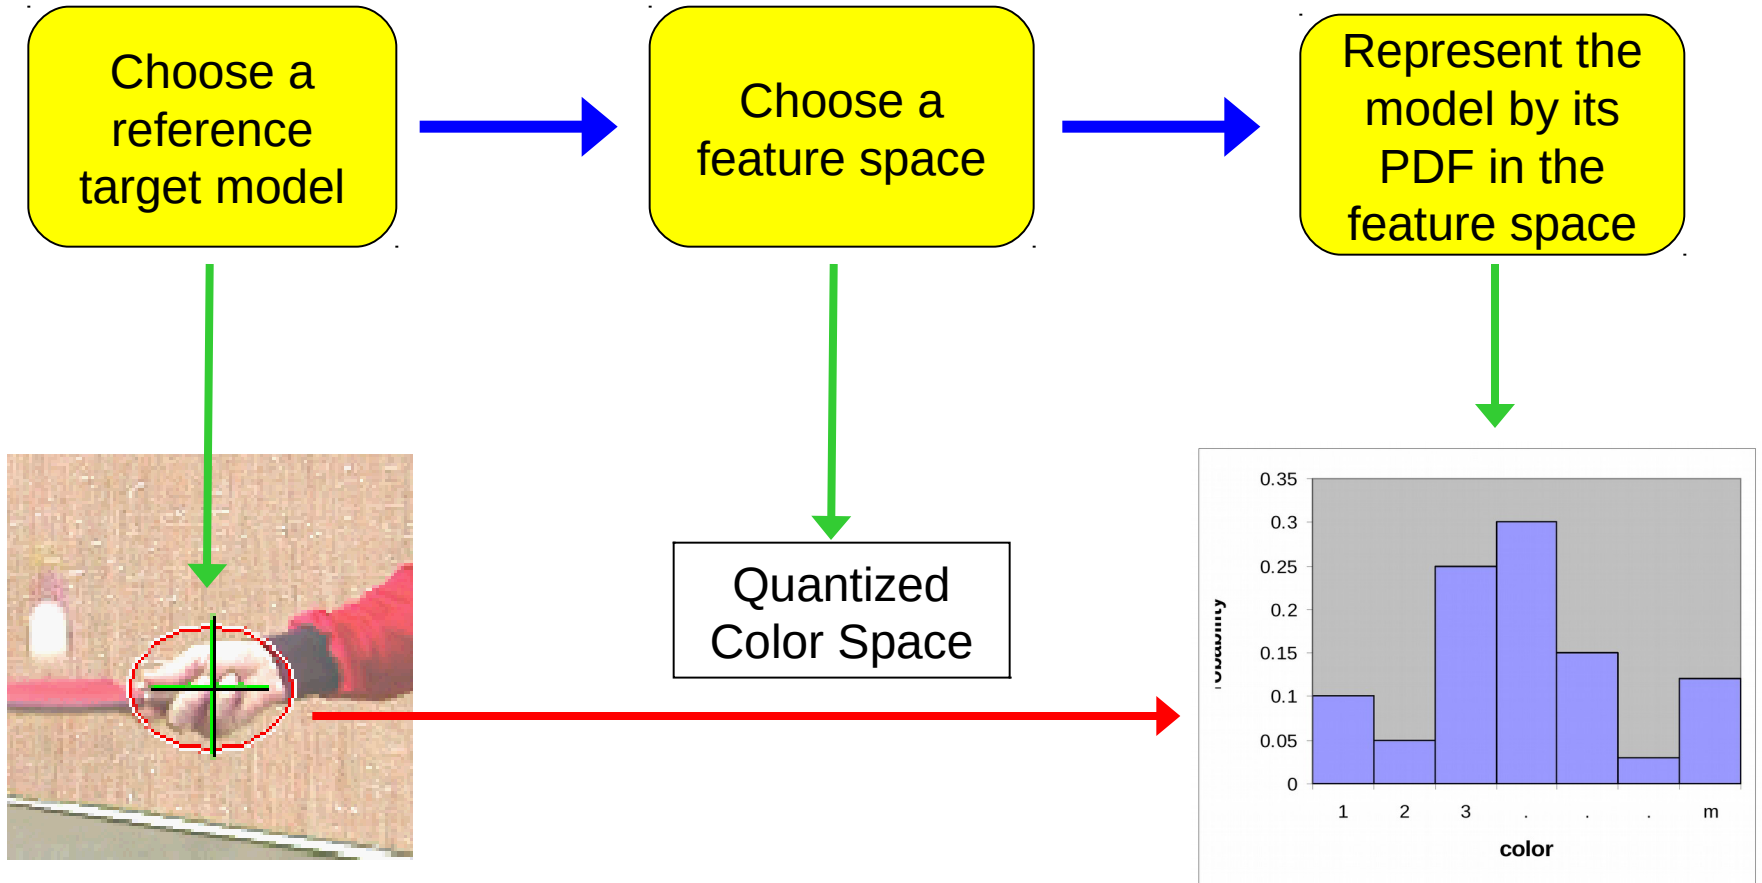
Mean-Shift Object Tracking

General Framework: Target Localization



Mean-Shift Object Tracking

Target Representation




Use Swain & Ballard's histogram backprojection

- Provides the PDF
- Do backprojection in the neighbourhood of the object location in previous frame
- Use mean shift to find where the mode (peak) is in the current frame

Mean-Shift Object Tracking

Maximizing the Similarity Function


Important Assumption:



The target representation provides sufficient discrimination



One mode in the searched neighborhood





End of CMPT 412 Notes

- Other details of improved tracking are contained in following slides.


Mean-Shift Object Tracking

PDF Representation

Target Model
(centered at 0)



Target Candidate
(centered at y)



$$\mathbf{r}_q = \{q_u\}_{u=1..m} \quad \sum_{u=1}^m q_u = 1$$

$$\mathbf{r}_{p(y)} = \{p_u(y)\}_{u=1..m} \quad \sum_{u=1}^m p_u = 1$$

Similarity
Function:


$$f(y) = f[\mathbf{r}_q, \mathbf{r}_{p(y)}]$$

Mean-Shift Object Tracking


Finding the PDF of the target model

$$\{x_i\}_{i=1..n}$$

Target pixel locations



0



y

$$k(x)$$

A differentiable, isotropic, convex, monotonically decreasing kernel

- Peripheral pixels are affected by occlusion and background interference


$$b(x)$$

The color bin index (1..m) of pixel x

Probability of feature u in model

$$q_u = C \sum_{b(x_i)=u} k(\|x_i\|^2)$$

Normalization factor




Pixel weight

Probability of feature u in candidate

$$p_u(y) = C_h \sum_{b(x_i)=u} k\left(\left\|\frac{y - x_i}{h}\right\|^2\right)$$

Normalization factor



Pixel weight

Mean-Shift Object Tracking

Similarity Function

Target model: $\mathbf{r}_q = (q_1, \mathbf{K}, q_m)$


Target candidate: $\mathbf{r}_{p(y)} = (p_1(y), \mathbf{K}, p_m(y))$

Similarity function: $f(y) = f[\mathbf{r}_{p(y)}, \mathbf{r}_q] = ?$

The Bhattacharyya Coefficient

$$\mathbf{r}'_q = (\sqrt{q_1}, \mathbf{K}, \sqrt{q_m})$$

$$\mathbf{r}'_{p(y)} = (\sqrt{p_1(y)}, \mathbf{K}, \sqrt{p_m(y)})$$



$$f(y) = \cos \theta_y = \frac{p'(y)^T q'}{\|p'(y)\| \|q'\|} = \sum_{u=1}^m \sqrt{p_u(y) q_u}$$


Mean-Shift Object Tracking

Target Localization Algorithm

Start from the position of the model in the current frame


Search in the model's neighborhood in next frame

Find best candidate by maximizing a similarity func.



q

$p(y)$



$$f[p(y), q]$$

Mean-Shift Object Tracking

Approximating the Similarity Function

$$f(y) = \sum_{u=1}^m \sqrt{p_u(y) q_u}$$

Model location: y_0

Candidate location: y

Linear
approx.
(around y_0)

$$f(y) \approx \underbrace{\frac{1}{2} \sum_{u=1}^m \sqrt{p_u(y_0) q_u}}_{\text{Independent of } y} + \frac{1}{2} \sum_{u=1}^m p_u(y) \sqrt{\frac{q_u}{p_u(y_0)}}$$

Independent
of y

$$p_u(y) = C_h \sum_{b(x_i)=u} k\left(\left\|\frac{y-x_i}{h}\right\|^2\right)$$

$$\frac{C_h}{2} \sum_{i=1}^n w_i k\left(\left\|\frac{y-x_i}{h}\right\|^2\right)$$


Density
estimate!
(as a function
of y)

Mean-Shift Object Tracking

Maximizing the Similarity Function

The mode of $\frac{C_h}{2} \sum_{i=1}^n w_i k \left(\left\| \frac{y - x_i}{h} \right\|^2 \right) = \text{sought maximum}$


Important Assumption:



The target representation provides sufficient discrimination



One mode in the searched neighborhood



Mean-Shift Object Tracking

Applying Mean-Shift

The mode of $\frac{C_h}{2} \sum_{i=1}^n w_i k \left(\left\| \frac{y - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}$ = sought maximum

Original
Mean-Shift:

Find mode of $c \sum_{i=1}^n k \left(\left\| \frac{y - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}$ using

$$y_1 = \frac{\sum_{i=1}^n x_i g \left(\left\| \frac{y_0 - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}}{\sum_{i=1}^n g \left(\left\| \frac{y_0 - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}}$$

Extended
Mean-Shift:

Find mode of $c \sum_{i=1}^n w_i k \left(\left\| \frac{y - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}$ using

$$y_1 = \frac{\sum_{i=1}^n x_i w_i g \left(\left\| \frac{y_0 - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}}{\sum_{i=1}^n w_i g \left(\left\| \frac{y_0 - x_i}{h} \right\|^2 \right)^{\frac{1}{2}}}$$

Mean-Shift Object Tracking

About Kernels and Profiles

A special class of radially symmetric kernels:

$$K(x) = ck(\|x\|^2)$$

The profile of kernel K

$$k'(x) = -g(x)$$

Extended
Mean-Shift:

Find mode of

$$c \sum_{i=1}^n w_i k\left(\left\|\frac{y - x_i}{h}\right\|^2\right)$$

using

$$y_1 = \frac{\sum_{i=1}^n x_i w_i g\left(\left\|\frac{y_0 - x_i}{h}\right\|^2\right)}{\sum_{i=1}^n w_i g\left(\left\|\frac{y_0 - x_i}{h}\right\|^2\right)}$$


Mean-Shift Object Tracking

Choosing the Kernel

A special class of radially symmetric kernels:


$$K(x) = ck(\|x\|^2)$$

Epanechnikov kernel:



$$k(x) = \begin{cases} 1 - \|x\|^2 & \text{if } \|x\| \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

Uniform kernel:




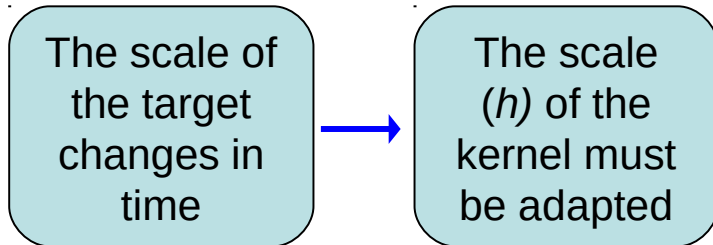
$$g(x) = -k(x) = \begin{cases} 1 - \|x\|^2 & \text{if } \|x\| \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

$$y_1 = \frac{\sum_{i=1}^n x_i w_i g\left(\left\|\frac{y_0 - x_i}{h}\right\|^2\right)}{\sum_{i=1}^n w_i g\left(\left\|\frac{y_0 - x_i}{h}\right\|^2\right)} \longrightarrow y_1 = \frac{\sum_{i=1}^n x_i w_i}{\sum_{i=1}^n w_i}$$



Mean-Shift Object Tracking

Adaptive Scale

Problem:




Solution:



Mean-Shift Object Tracking

Results



Feature space: $16 \times 16 \times 16$ quantized RGB

Target: manually selected on 1st frame

Average mean-shift iterations: 4

Mean-Shift Object Tracking

Results




Partial occlusion



Distraction




Motion blur



Mean-Shift Object Tracking

Results



Mean-Shift Object Tracking

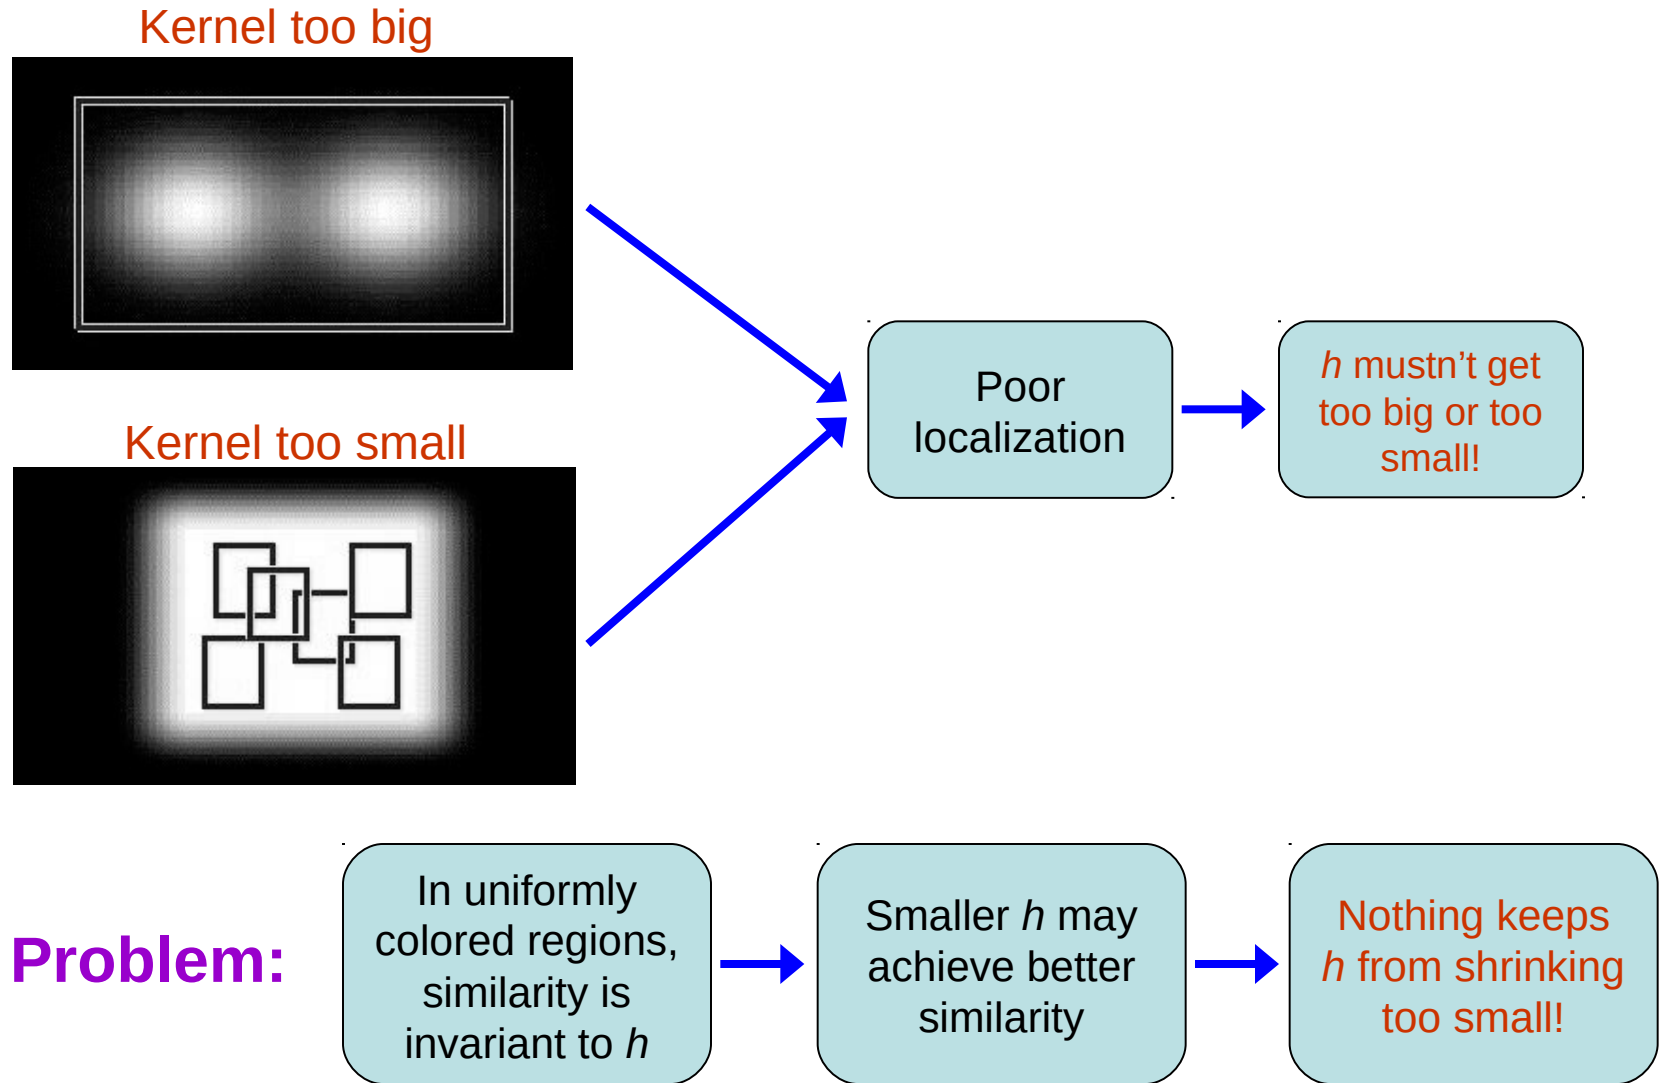
Results



Feature space: 128×128 quantized RG

Mean-Shift Object Tracking

The Scale Selection Problem



Tracking Through Scale Space

Motivation



Spatial
localization
for several
scales

Previous method





Simultaneous
localization in
space and
scale

This method

Lindeberg's Theory

Selecting the best scale for describing image features



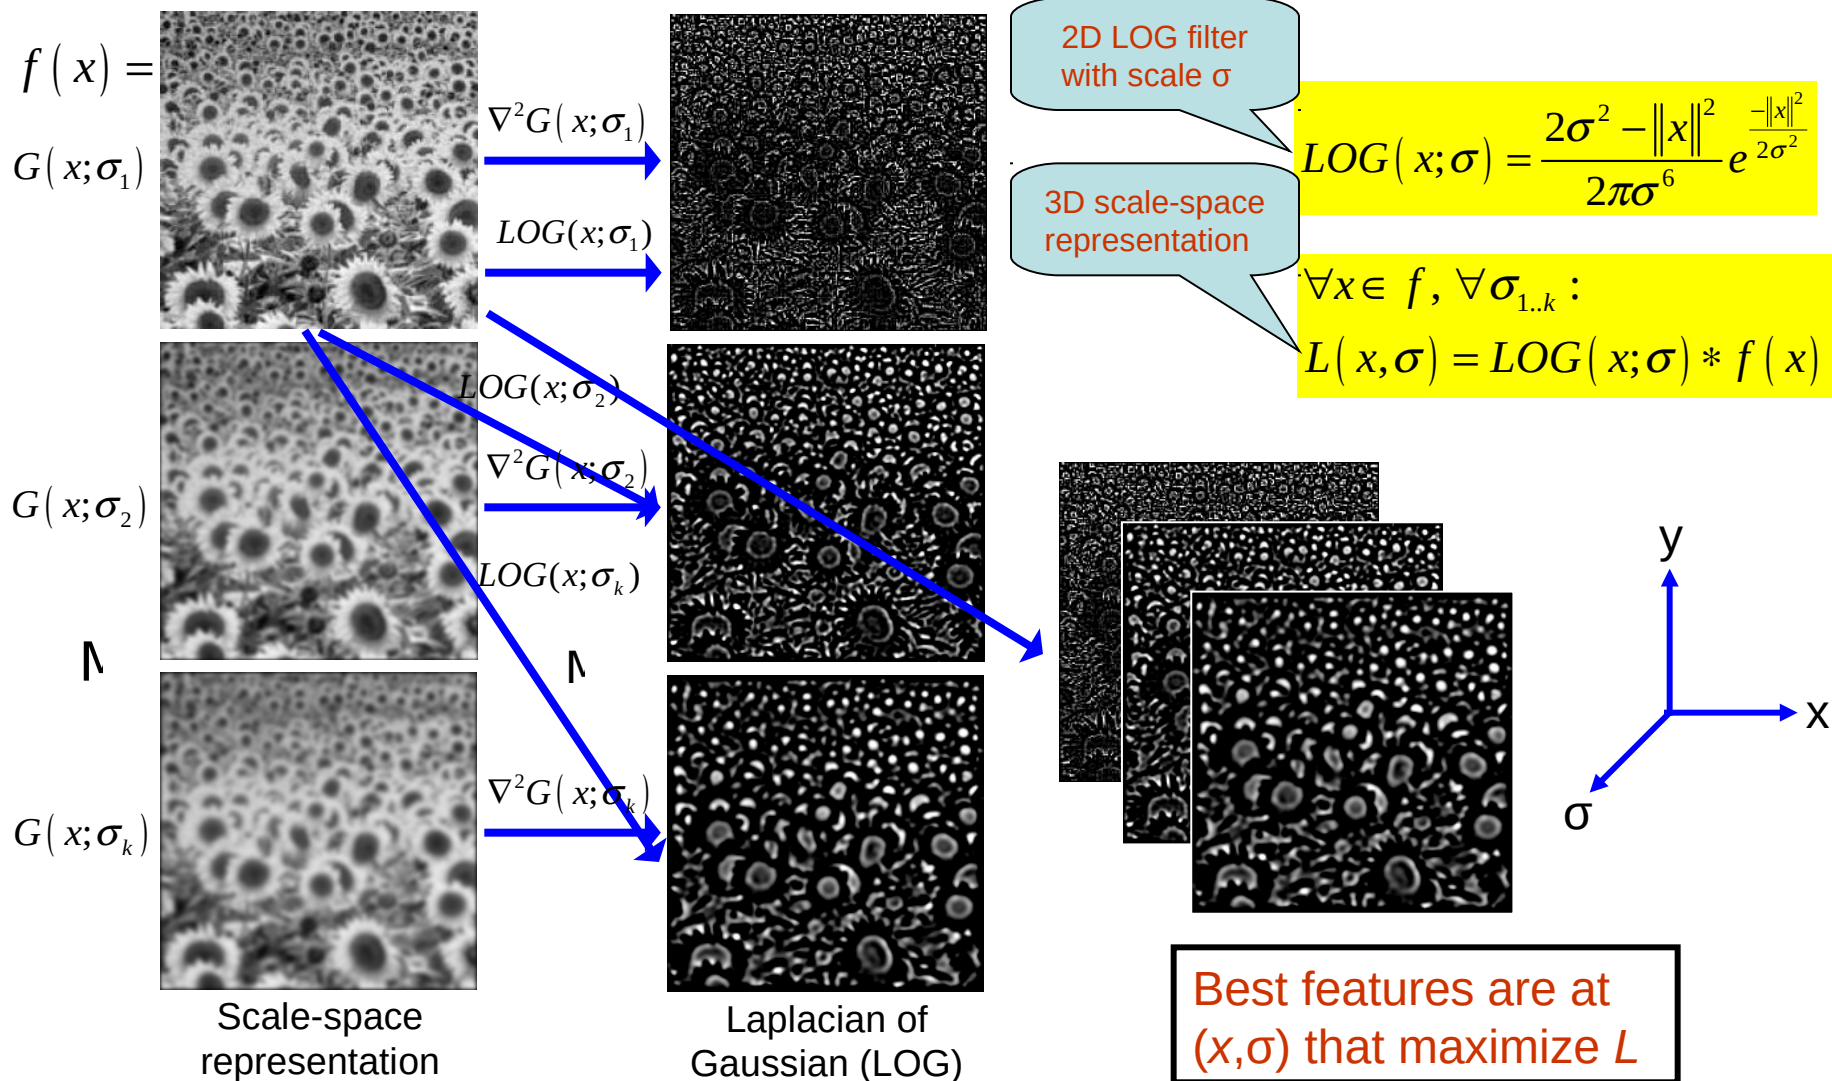
Scale-space
representation

Differential
operator applied

50 strongest
responses


Lindeberg's Theory

The Laplacian operator for selecting blob-like features



Lindeberg's Theory

Multi-Scale Feature Selection Process



Tracking Through Scale Space

Approximating LOG using DOG

$$LOG(x; \sigma) \approx DOG(x; \sigma) = G(x; \sigma) - G(x; 1.6\sigma)$$

2D LOG filter
with scale σ

2D DOG filter
with scale σ

2D Gaussian
with $\mu=0$ and
scale σ


2D Gaussian
with $\mu=0$ and
scale 1.6σ

Why DOG?

- Gaussian pyramids are created faster
- Gaussian can be used as a mean-shift kernel

3D spatial
kernel

$$K(x, \sigma) =$$



DOG filters at
multiple scales



Scale-space
filter bank

Tracking Through Scale Space

Using Lindeberg's Theory



Weight image



$$= E(x, \sigma)$$

3D scale-space representation

Recall:

Model: $\mathbf{r} \quad \mathbf{q} = (q_1, \mathbf{K}, q_m)$ at y_0

Candidate: $\mathbf{r} \quad \mathbf{p}(y) = (p_1(y), \mathbf{K}, p_m(y))$

Color bin: $b(x)$

Pixel weight: $w(x) = \sqrt{\frac{q_{b(x)}}{p_{b(x)}(y_0)}}$

Centered at
current
location and
scale


The likelihood
that each
candidate
pixel belongs
to the target

Modes are blobs in
the scale-space
neighborhood

Need a mean-shift
procedure that
finds local modes
in $E(x, \sigma)$

Tracking Through Scale Space

Example



Tracking Through Scale Space

Applying Mean-Shift


Use interleaved spatial/scale mean-shift

Spatial stage:

Fix σ and
look for the
best x

Scale stage:

Fix x and
look for the
best σ



Iterate stages
until
convergence of
 x and σ


Tracking Through Scale Space

Results

Fixed-scale



$\pm 10\%$ scale adaptation



Tracking through scale space



Thank
You